

1 A brief overview of the effect-system based encoding of the session-typed $\pi + \lambda$ -calculus into Haskell

The basis of the $\pi + \lambda$ encoding in Haskell is a *graded monad* which is used to track session information. This is encoded via the data type:

```
data Session (s :: [*]) a = Session { getProcess :: IO a }
```

This wraps the *IO* monad in a binary type constructor *Session* with deconstructor *getProcess* :: *Session s a* → *IO a* and with a tag *s* used for type-level session information. In practise, we only need *getProcess* internally, so this can be hidden. We define a type-refined version of *getProcess* which allows us to run a computation only when the session environment is empty, that is, the process is closed with respect to channels.

```
run :: Session '[] a → IO a
run = getProcess
```

We can therefore run any session which will evaluate everything inside of the *IO* monad and actually performing the communication/spawning/etc.

Type-level session information will take the form of a list of mappings from channel names to session types, written like: `'[c :→ S, d :→ T, ...]`. This list will get treated as a set when we compose computations together, that is there are no duplicate mappings of some channel variable *c*, and the ordering will be normalise (this is a minor point and shouldn't affect too much here).

Session types are defined by the following type constructors:

```
-- Session types
data a !: s
data a :? s
data End
```

Duality of session type is then defined as a simple type-level function:

```
type family Dual s where
  Dual End = End
  Dual (t !: s) = t :? (Dual s)
  Dual (t :? s) = t !: (Dual s)
```

We define a (finite) set of channel name symbols *ChanNameSymbol* [this can be generalised away, but for some slightly subtle reasons mostly to do with CloudHaskell internals I have avoided the generalisation for the moment].

```
data ChanNameSymbol = X | Y | Z | C | D
data ChanName = Ch ChanNameSymbol | Op ChanNameSymbol
```

ChanName thus can describe the dual end of a channel with *Op*. These are just names for channels. Channels themselves comprise an encapsulated Concurrent Haskell channel [todo: convert to a Cloud Haskell channel]

```
data Channel (n :: ChanName) = forall a ◦ Channel (C.Chan a) deriving Typeable
```

1.1 π -calculus part

We can now define the core primitives for send and receive, which have types:

```
send :: Channel c → t → Session '[c :→ t :! End] ()
```

```
recv :: Channel c → Session '[c :→ t :? End] t
```

These both take a named channel *Channel c* and return a *Session* computation indexed by the session environment $'[c :→ S]$ where *S* is either a send or receive action (terminated by *End*). These can then be composed using the **do**-notation, which sequentially composes session information. For example:

```
data Ping = Ping deriving Show
data Pong = Pong deriving Show
foo (c :: Channel (Ch C)) = do send c Ping
                             x ← recv c
                             return ((x + 1) :: Int)
```

This function is of type:

```
foo :: Channel (Ch C) → Session '[Ch C :→ (Ping :! (Int :? End))] Int
```

describing the session channel behaviour for *C*.

I've given an explicit name to the channel *c* here via a type signature, which names it as *Ch C*. This isn't strictly necessary here, but it leads to a huge simplification in the inferred type.

The *new* combinator then models ν , which takes a function mapping from a pair of two channels names *Ch c* and *Op C* to a session with behaviour *s*, and creates a session where any mention to *Ch c* or *Op c* is removed:

```
new :: (Duality s c) ⇒
      ((Channel (Ch c), Channel (Op c)) → Session s b)
      → Session (Del (Ch c) (Del (Op c) s)) b
```

That is, the channels *Ch c* and *Op c* are only in scope for *Session s b*.

The *Duality* predicate asks whether the session environment *s* contains dual session types for channel *Ch C* and its dual *Op c*.

The effect monad used to sequence session information also has some predicates to check that dual session channels do not appear in the same (sequential) session. For example, the following is a type error:

```

% This code fails to type check
foo2a = new ( $\lambda(c :: (\text{Channel } (\text{Ch } C)), c' :: (\text{Channel } (\text{Op } C))) \rightarrow$ 
  do Ping  $\leftarrow$  recv c'
  send c Ping
  return ())

```

The produces a type error:

```

cloudh.lhs:283:23:
  Couldn't match type 'DualConflict ('Op 'C)' with 'NoSeqDuals'
  In a stmt of a 'do' block: Ping <- recv c'
  .....

```

(In the case of delegation, this check is not used, so delegated channels can appear in sequential composition with their dual channel names- this is an interesting subtlety that requires a bit of extra work in the background. It is not seen at the top-level though).

To use channels properly, we need parallel composition. This is given by:

$$\text{par} :: (\text{Disjoint } s \ t) \Rightarrow \text{Session } s \ () \rightarrow \text{Session } t \ () \rightarrow \text{Session } (\text{UnionS } s \ t) \ ()$$

The binary predicate *Disjoint* here checks that *s* and *t* do not contain any of the same channels. *UnionS* takes the disjoint union of the two environments.

We can now define a complete example with communication:

```

server c = do Ping <- recv c
           print "Server: Got a ping"
process = new ( $\lambda(c, c') \rightarrow \text{par } (\text{send } c \text{ Ping}) (\text{server } c')$ )

```

Which we can run with *run process* getting "Server: Got a ping". Note that the types here are completely inferred, giving *process* :: *Session* '[] ()

1.1.1 Delegation

So far we have dealt with only first-order channels (in the sense that they can pass only values and not other channels). We introduce a “delegate” type to wrap the session types of channels being passed:

```
data DelgS s
```

Channels can then be sent with *chSend* primitive:

$$\text{chSend} :: \text{Channel } c \rightarrow \text{Channel } d \rightarrow \text{Session } '[c \rightarrow (\text{DelgS } s) \text{ :! End}, d \rightsquigarrow s] \ ()$$

i.e., we can send a channel *d* with session type *s* over *c*.

[Note, the different mapping \rightsquigarrow instead of \rightarrow for *d*, this marks to the type system that this channel has been delegated, therefore it can appear in the same environment as the dual channel of *d*.]

The dual of this is a little more subtle. Receiving a delegated channel is given by combinator, which is not a straightforward monadic function, but takes a function as an argument:

$$\begin{aligned} \text{chRecv} :: \text{Channel } c \rightarrow (\text{Channel } d \rightarrow \text{Session } s \ a) \rightarrow \\ \text{Session } (\text{UnionS } '[c \rightarrow (\text{DelgS } (\text{Lookup } s \ d)) :? (\text{Lookup } s \ c)] (\text{Del } d \ s)) \ a \end{aligned}$$

Given a channel c , and a computation which binds channel d to produces behaviour c , then this is provided by receiving d over c . Thus the resulting computation is the union of c mapping to the session type of d in the session environment s , composed with the s but with d deleted (removed).

Here is an example using delegation. Consider the following process *server2* which receives a channel d on c , and then sends a ping on it:

$$\begin{aligned} \text{server2 } c = \text{chRecv } c \\ (\lambda(d :: \text{Channel } (\text{Ch } D)) \rightarrow \text{send } d \text{ Ping}) \end{aligned}$$

(Note, I have had to include explicit types to give a concrete name to the channel d , this is an unfortunate artefact of the current encoding, but not too bad from a theoretical perspective).

The type of *server2* is inferred as:

$$\text{server2} :: \text{Channel } c \rightarrow \text{Session } '[c \rightarrow (\text{DelgS } (\text{Ping} :! \text{End}) :? \text{Lookup } '[\text{'Ch } ' D \rightarrow (\text{Ping} :! \text{End})] \ c)] \ ()$$

We then define a client to interact with this that binds d (and its dual d'), then sends d over c and waits to receive a ping on d'

$$\begin{aligned} \text{client2 } (c :: \text{Channel } (\text{Ch } C)) = \\ \text{new } (\lambda(d :: (\text{Channel } (\text{Ch } D)), d') \rightarrow \\ \text{do } \text{chSend } c \ d \\ \text{Ping} \leftarrow \text{recv } d' \\ \text{print "Client: got a ping"}) \end{aligned}$$

This has inferred type:

$$\text{client2} :: \text{Dual } s \sim (\text{Ping} :? \text{End}) \Rightarrow \text{Channel } (\text{'Ch } ' C) \rightarrow \text{Session } '[\text{'Ch } ' C \rightarrow (\text{DelgS } s :! \text{End})] \ ()$$

The type constraint says that the dual of s is a session that receives a *Ping*, so s is *Ping* :! *End*.

We then compose *server2* and *client2* in parallel, binding the channels c and its dual c' to give to client and server.

$$\text{process2} = \text{new } (\lambda(c, c') \rightarrow \text{par } (\text{client2 } c) (\text{server2 } c'))$$

This type checks and can be then run (*run process2*) yielding "Client: got a ping".

1.2 λ -part

Since we are studying the $\pi + \lambda$ -calculus, we can abstract over channels with linear functions. So far we have abstracted over channels, but not in an *operational sense*- think of this more as let-binding style substitution (cut). We now introduce linear functions which can abstract over channels (and the session types of those channels, which the previous form of abstraction above **doesn't do**, it just abstracts over names, not the session types associated with their names).

First, we abstract functions via a type constructor *Abs*

data *Abs* $t\ s = \text{Abs} (\text{Proxy } s) (\text{forall } c \circ (\text{Channel } c \rightarrow \text{Session } (\text{UnionS } s \text{ '}[c:\rightarrow t]\text{') } ()))$

The *Abs* data constructor takes a function of type $\text{forall } c \circ (\text{Channel } c \rightarrow \text{Session } (\text{UnionS } s \text{ '}[c:\rightarrow t]\text{') } ())$, that is, a function from *universally quantified* channel name c to a *Session* environment s where $c:\rightarrow t$ is a member). Since *UnionS* is a non-injective function we also need a (trivial) type annotation that explains exactly what is the remaining channel- this is *Proxy* s (I'll shown an example in the moment). This returns a result *Abs* $t\ s$ which describes a function which takes some channel with session type t and has session environment s , cf.

$$\frac{\Delta, c : T \vdash C : \Diamond}{\Delta \vdash \lambda c. C : T \multimap \Diamond}$$

This can then be applied by the following primitive

appH :: *Abs* $t\ s \rightarrow \text{Channel } c \rightarrow \text{Session } (\text{UnionS } s \text{ '}[c:\rightarrow t]\text{') } ()$
appH (*Abs* $_k$) $c = k\ c$

Thus, given a linear session function *Abs* $t\ s$ and some channel c then we get a session with environment s and a mapping $c:\rightarrow t$. Here's an example: a client abstract over a channel, and then applies it within the same process:

client4 ($c :: \text{Channel } (\text{Ch } C)$) = **do**
let $f = \text{Abs} (\text{Proxy} :: (\text{Proxy } '[])) (\lambda c \rightarrow \text{send } c\ \text{Ping})$
appH $f\ c$

This simply has type *client4* :: $\text{Channel } (' \text{Ch } ' C) \rightarrow \text{Session } ' [' \text{Ch } ' C : \rightarrow (\text{Ping} :! \text{End})] ()$. We can then interact with this in a usual straightforwad way.

process4 = *new* ($\lambda(c, c') \rightarrow (\text{client4 } c) \text{ 'par' } (\text{do } \{ x \leftarrow \text{recv } c'; \text{print } x \})$)