

1 A brief overview of the effect-system based encoding of the session-typed $\pi + \lambda$ -calculus into Haskell

The basis of the $\pi + \lambda$ encoding in Haskell is a *graded monad* which is used to track session information. This is encoded via the data type:

```
data Session (s :: [*]) a = Session {getProcess :: IO a}
```

This wraps the *IO* monad in a binary type constructor *Session* with deconstructor *getProcess* :: *Session s a* → *IO a* and with a tag *s* used for type-level session information. In practise, we only need *getProcess* internally, so this can be hidden. We define a type-refined version of *getProcess* which allows us to run a computation only when the session environment is empty, that is, the process is closed with respect to channels.

```
run :: Session '[] a → IO a
run = getProcess
```

We can therefore run any session which will evaluate everything inside of the *IO* monad and actually performing the communication/spawning/etc.

Type-level session information will take the form of a list of mappings from channel names to session types, written like: $[c \rightarrow S, d \rightarrow T, \dots]$. This list will get treated as a set when we compose computations together, that is there are no duplicate mappings of some channel variable *c*, and the ordering will be normalise (this is a minor point and shouldn't affect too much here).

Session types are defined by the following type constructors:

```
-- Session types
data a ! s
data a ? s
data End
```

Duality of session type is then defined as a simple type-level function:

```
type family Dual s where
  Dual End = End
  Dual (t ! s) = t ? (Dual s)
  Dual (t ? s) = t ! (Dual s)
  Dual (Sel l s t) = (Dual s) : & (Dual t)
  Dual (s1 : & s2) = Sel Sup (Dual s1) (Dual s2)
```

We define a (finite) set of channel name symbols *ChanNameSymbol* [this can be generalised away, but for some slightly subtle reasons mostly to do with CloudHaskell internals I have avoided the generalisation for the moment].

```
data ChanNameSymbol = X | Y | Z | C | D | ForAll -- reserved
data ChanName = Ch ChanNameSymbol | Op ChanNameSymbol
```

ChanName thus can describe the dual end of a channel with *Op*. These are just names for channels. Channels themselves comprise an encapsulated Concurrent Haskell channel [todo: convert to a Cloud Haskell channel]

data *Channel* (*n* :: *ChanName*) = forall *a* ◦ *Channel* (*C.Chan a*) **deriving** *Typeable*

1.1 π -calculus part

We can now define the core primitives for send and receive, which have types:

send :: *Channel c* → *t* → *Session* '[*c* :→ *t* ! End] ()

recv :: *Channel c* → *Session* '[*c* :→ *t* ? End] *t*

These both take a named channel *Channel c* and return a *Session* computation indexed by the session environment '[*c* :→ *S*]' where *S* is either a send or receive action (terminated by *End*). These can then be composed using the **do**-notation, which sequentially composes session information. For example:

```
data Ping = Ping deriving Show
data Pong = Pong deriving Show
foo (c :: Channel (Ch C)) = do send c Ping
                               x ← recv c
                               return ((x + 1) :: Int)
```

This function is of type:

foo :: *Channel* (*Ch C*) → *Session* '[*Ch C* :→ (*Ping* ! (*Int* ? End))] *Int*

describing the session channel behaviour for *C*.

I've given an explicit name to the channel *c* here via a type signature, which names it as *Ch C*. This isn't strictly necessary here, but it leads to a huge simplification in the inferred type.

The *new* combinator then models ν , which takes a function mapping from a pair of two channels names *Ch c* and *Op C* to a session with behaviour *s*, and creates a session where any mention to *Ch c* or *Op c* is removed:

```
new :: (Duality s c) ⇒
      ((Channel (Ch c), Channel (Op c)) → Session s b)
      → Session (Del (Ch c) (Del (Op c) s)) b
```

That is, the channels *Ch c* and *Op c* are only in scope for *Session s b*.

The *Duality* predicate asks whether the session environment *s* contains dual session types for channel *Ch C* and its dual *Op c*.

The session type encoding here is for an asynchronous calculus. In which case, the following is allowed:

```

foo2 = new (λ(c :: (Channel (Ch C)), c' :: (Channel (Op C))) →
  do Ping ← recv c'
  send c Ping
  return ())

```

To use channels properly, we need parallel composition. This is given by:

$$par :: (Disjoint\ s\ t) \Rightarrow Session\ s\ () \rightarrow Session\ t\ () \rightarrow Session\ (UnionS\ s\ t)\ ()$$

The binary predicate *Disjoint* here checks that *s* and *t* do not contain any of the same channels. *UnionS* takes the disjoint union of the two environments.

We can now define a complete example with communication:

```

server c = do Ping ← recv c
            print "Server: Got a ping"
process = new (λ(c, c') → par (send c Ping) (server c'))

```

Which we can run with *run process* getting "Server: Got a ping". Note that the types here are completely inferred, giving *process* :: *Session* '[]' ().

1.1.1 Delegation

So far we have dealt with only first-order channels (in the sense that they can pass only values and not other channels). We introduce a “delegate” type to wrap the session types of channels being passed:

```
data DelgS s
```

Channels can then be sent with *chSend* primitive:

$$chSend :: Channel\ c \rightarrow Channel\ d \rightarrow Session\ '[c : \rightarrow (DelgS\ s) :! End, d : \rightarrow s]\ ()$$

i.e., we can send a channel *d* with session type *s* over *c*.

The dual of this is a little more subtle. Receiving a delegated channel is given by combinator, which is not a straightforward monadic function, but takes a function as an argument:

$$chRecv :: Channel\ c \rightarrow (Channel\ d \rightarrow Session\ s\ a) \rightarrow Session\ (UnionS\ '[c : \rightarrow (DelgS\ (Lookup\ s\ d)) :? (Lookup\ s\ c)] (Del\ d\ s))\ a$$

Given a channel *c*, and a computation which binds channel *d* to produces behaviour *c*, then this is provided by receiving *d* over *c*. Thus the resulting computation is the union of *c* mapping to the session type of *d* in the session environment *s*, composed with the *s* but with *d* deleted (removed).

Here is an example using delegation. Consider the following process *server2* which receives a channel *d* on *c*, and then sends a ping on it:

```

server2 c = chRecv c
            (λ(d :: Channel (Ch D)) → send d Ping)

```

(Note, I have had to include explicit types to give a concrete name to the channel d , this is an unfortunate artefact of the current encoding, but not too bad from a theoretical perspect).

The type of *server2* is inferred as:

$server2 :: Channel\ c \rightarrow Session\ '[c:\rightarrow(DelgS\ (Ping:!End):?Lookup\ '[\text{'Ch'}\ D:\rightarrow(Ping:!End)]\ c)]\ ()$

We then define a client to interact with this that binds d (and its dual d'), then sends d over c and waits to receive a ping on d'

```
client2 (c :: Channel (Ch C)) =
  new (\(d :: (Channel (Ch D)), d') →
    do chSend c d
      Ping ← recv d'
      print "Client: got a ping")
```

This has inferred type:

$client2 :: Dual\ s \sim (Ping:?End) \Rightarrow Channel\ (\text{'Ch'}\ C) \rightarrow Session\ '[\text{'Ch'}\ C:\rightarrow(DelgS\ s:!End)]\ ()$

The type constraint says that the dual of s is a session that receives a *Ping*, so s is *Ping:!End*.

We then compose *server2* and *client2* in parallel, binding the channels c and its dual c' to give to client and server.

$process2 = new\ (\lambda(c, c') \rightarrow par\ (client2\ c)\ (server2\ c'))$

This type checks and can be then run (*run process2*) yielding "Client: got a ping".

1.2 λ -part

Since we are studying the $\pi + \lambda$ -calculus, we can abstract over channels with linear functions. So far we have abstracted over channels, but not in an *operational sense*- think of this more as let-binding style substitution (cut). We now introduce linear functions which can abstract over channels (and the session types of those channels, which the previous form of abstraction above **doesn't do**, it just abstracts over names, not the session types associated with their names).

First, we abstract functions via a type constructor *Abs*

data *Abs* $t\ s = forall\ c \circ Abs\ (Proxy\ s)\ (Channel\ c \rightarrow Session\ (UnionS\ s\ '[c:\rightarrow t])\ ())$

The *Abs* data constructor takes a function of type $(Channel\ c \rightarrow Session\ (UnionS\ s\ '[c:\rightarrow t])\ ())$, that is, a function from some channel c to a *Session* environment s where $c:\rightarrow t$ is a member). Since *UnionS* is a non-injective function we also need a type annotation that explains exactly what is the remaining channel- this is *Proxy s* (I'll show an example in the moment). This returns a result *Abs t s* which describes a function which takes some channel with session type t and has session environment s , cf.

$$\frac{\Delta, c : T \vdash C : \diamond}{\Delta \vdash \lambda c. C : T \multimap \diamond}$$

This can then be applied by the following primitive

$$\begin{aligned} appH &:: Abs\ t\ s \rightarrow Channel\ c \rightarrow Session\ (UnionS\ s\ ['c \rightarrow t])\ () \\ appH\ (Abs\ _k)\ c &= \mathbf{let}\ (Session\ s) = k\ (unsafeCoerce\ c)\ \mathbf{in}\ Session\ s \end{aligned}$$

Whatever concrete name was used for the channel in the abstracted process is replaced by the channel name here.

Thus, given a linear session function $Abs\ t\ s$ and some channel c then we get a session with environment s and a mapping $c \mapsto t$. Here's an example: a client abstract over a channel, and then applies it within the same process:

$$\begin{aligned} & client_4 \ (c :: Channel \ (Ch \ C)) = \mathbf{do} \\ & \quad \mathbf{let} \ f = Abs \ (Proxy :: (Proxy \ '[])) \ (\lambda c \rightarrow send \ c \ Ping) \\ & \quad appH \ f \ c \end{aligned}$$

This simply has type $client_4 :: Channel \text{ 'Ch ' } C \rightarrow Session \text{ ' ['Ch ' } C : \rightarrow (Ping :! End)] ()$. We can then interact with this in a usual straightforward way.

$$process_4 = new (\lambda(c, c') \rightarrow (client_4\ c) \text{ 'par' } (\mathbf{do} \{ x \leftarrow recv\ c'; print\ x \}))$$

A more complicated example reuses the abstraction in the client with different channels:

$$\begin{aligned} & client5 \ (c :: Channel \ (Ch \ C)) \ (d :: (Channel \ (Ch \ D))) \ (x :: (Channel \ (Ch \ X))) = \mathbf{do} \\ & \quad \mathbf{let} \ f = Abs \ (Proxy :: (Proxy \ '[(Ch \ X) \rightarrow Pong \ ! \ End])) \\ & \quad (\lambda(c :: (Channel \ (Ch \ D))) \rightarrow \mathbf{do} \ send \ c \ Ping \\ & \quad \quad send \ x \ Pong) \\ & \quad appH \ f \ c \\ & \quad appH \ f \ d \end{aligned}$$
$$\begin{aligned}
& process5 = new (\lambda(c :: Channel (Ch C), c') \rightarrow \\
& \quad new (\lambda(x :: Channel (Ch X), x') \rightarrow \\
& \quad \quad new (\lambda(d :: Channel (Ch D), d') \rightarrow \\
& \quad \quad \quad (client5\ c\ d\ x)\ 'par' \\
& \quad \quad \quad \mathbf{do}\ v \leftarrow recv\ c' \\
& \quad \quad \quad \quad print\ v \\
& \quad \quad \quad \quad v \leftarrow recv\ x' \\
& \quad \quad \quad \quad print\ v \\
& \quad \quad \quad \quad v \leftarrow recv\ d' \\
& \quad \quad \quad \quad print\ v \\
& \quad \quad \quad \quad v \leftarrow recv\ x' \\
& \quad \quad \quad \quad print\ v \\
& \quad \quad \quad \quad)))
\end{aligned}$$

Dimtris example

```

client6 (c :: (Channel (Ch C))) (d :: (Channel (Ch D))) =
  do let f = Abs (Proxy :: Proxy '[(Ch C) :→ Int ! End])
    (λ(z :: (Channel (Ch Z))) → (send z 42 'par' send c 7))
    appH f d

```

```

process6 = new (λ(c :: (Channel (Ch C)), c') →
  new (λ(d :: (Channel (Ch D)), d') →
    do client6 c d
      v1 ← recv c'
      v2 ← recv d'
      return (v1 + v2)))

```

1.3 Branching and choice

To encode branching and choice, we introduce binary branch/select (from which more complicated branch/select can be encoded) with two labels:

```

data Left
data Right
data Sup

```

```

data Label l where
  LeftL :: String → Label Left
  RightL :: String → Label Right

```

The label data constructors *LeftL* and *RightL* also take string parameters for convenience (to act as comments in the code).

Note that whilst 'Sup' is a viable type-level label, there is no way to construct a label value with this type index. This is used for subtyping, where *Sup* represents a selection type which is a supertype.

Selection and branching session types are provided by the following two type constructors respectively:

```

data Sel l s t
data s : & t

```

Select then has the type:

```

select :: Channel c → Label l → Session '[c :→ Sel l End End] ()

```

The idea is that, given a channel *c*, and a label *l*, then a session is created with a select session type for label *l*. Any computations that get composed after that use *c* will add their session types into branch corresponding to the label. For example:

```

foo3 (c :: (Channel (Ch C))) =
  do select c (LeftL "1")
    v ← recv c
    send c (42 :: Int)

```

foo3 has the inferred type:

$$foo3 :: Channel \text{ 'Ch ' } C \rightarrow Session \text{ '['Ch ' } C : \rightarrow Sel \text{ Left } (t : ?(Int : !End)) \text{ End] } ()$$

That is, we see that after selecting the left branch, then *c* is used to receive some *t* and then send an *Int*.

Branching then has the following type:

$$\begin{aligned}
branch :: ((Del \text{ c } s1) \sim (Del \text{ c } s2)) \Rightarrow \\
& Channel \text{ c } \rightarrow (Label \text{ Left } \rightarrow Session \text{ s1 a}) \\
& \rightarrow (Label \text{ Right } \rightarrow Session \text{ s2 a}) \\
& \rightarrow Session (UnionS (Del \text{ c } s1) \text{ '[c : } \rightarrow ((Lookup \text{ s1 c}) : \& (Lookup \text{ s2 c}))]) a
\end{aligned}$$

This is a bit more complicated. The first parameter is the channel over which a choice is being offered. Then come two continuations, the process if the left branch is taken and the process if the right branch is taken. Each gives a session environment *s1* and *s2* but apart from a session type for *c*, these must be equal (shown by the constraint $(Del \text{ c } s1) \sim (Del \text{ c } s2)$). Finally, the returned session is that of $(Del \text{ c } s1)$ unioned with *c* mapping to the $(Lookup \text{ s1 c}) : \& (Lookup \text{ s2 c})$, i.e., the branching pair of the session types for *c* in the left and right branches.

Here's an example:

```

process7 = new (\(c :: (Channel (Ch C)), c') →
  do { select c (LeftL ""); send c 42 }
  'par' branch c' (\(LeftL "") → do { v ← recv c'; print v })
    (\(RightL "") → do { return (); return () })

```

Then *run process7* yields 42 as expected.

$$\begin{aligned}
selSupL :: Session \text{ '[c : } \rightarrow Sel \text{ l s End] } () \rightarrow Session \text{ '[c : } \rightarrow Sel \text{ Sup s t] } () \\
selSupL \text{ s} = Session \$ getProcess \text{ s}
\end{aligned}$$

$$\begin{aligned}
selSupR :: Session \text{ '[c : } \rightarrow Sel \text{ l End s] } () \rightarrow Session \text{ '[c : } \rightarrow Sel \text{ Sup t s] } () \\
selSupR \text{ s} = Session \$ getProcess \text{ s}
\end{aligned}$$

2 Hotel booking scenario

$$\begin{aligned}
p :: (?room :: String, ?credit :: Int) \Rightarrow \\
Channel (Ch X) \rightarrow Abs (Int : ! (End : \& End)) \text{ '[(Ch X) : } \rightarrow String : ! (Int : ? (Sel \text{ Sup } (Int : ! End) \text{ End})
\end{aligned}$$

```

p x = Abs (Proxy :: Proxy '[(Ch X) :→ String ! (Int :? (Sel Sup (Int ! End) End))])
(λ(y :: (Channel (Ch Y))) →
  do send x ? room
    quote ← recv x
    send y quote
    branch y (λ(LeftL "accept") → selSupL $ do select x (LeftL "accept")
              send x ? credit)
            (λ(RightL "reject") → selSupR $ select x (RightL "reject")))

```

```

{-foofoo (s1 :: (Channel (Ch Y))) (s2 :: (Channel (Ch Z))) (h1 :: (Channel (Ch C))) (h2 :: (Channel (

```